

Computer Vision & Pattern Recognition


Spring 2025

Assignment 7

May 9, 2025

Problem 1 [10 points]

Consider the pixel image “house1.png” (you can find it on iCorsi) and write a program for recovering the *internal* and *external* camera parameters of the camera that was used to generate this picture. In order to do so, you can find the *world coordinates* $X_i, i = 1, \dots, 10$, of 10 cardinal points of the house marked here



in the file “coords.tex” (you can find it on iCorsi). Following the lecture, you should attend to the following tasks:

1. Find the image coordinates $x_i, i = 1, \dots, 10$, of the points marked in the image above (either by detecting lines with the Hough transform and intersecting them or by reading them off the image “house1.png” manually).
2. Reconstruct the *projection matrix* P using the DLT algorithm. Do not worry if the rank of the matrix A that is involved is not 11. In practice, a slight imprecision in the coordinates is enough to turn it into a rank 12 matrix, but taking as p (the \mathbb{R}^{12} vector with the coefficients of P) the last column of V (not of V^T) of the singular value decomposition $A = U\Sigma V^T$ is still the way to go.
3. Recover the *camera calibration matrix* K , the *camera orientation* R and the (non-homogeneous) *world coordinates* \tilde{C} of the camera from P . Make sure that the diagonal values of K are positive in your solution by suitably adapting R .

Do the same also for the second image “house2.png” (you can find it on iCorsi).

Hand in your code, a short description of your solution, and your estimates of K , R , and \tilde{C} for both images.

Solutions must be returned on May 20, 2025 via iCorsi